

A Theory of Efficient Price Cycles

Under Imperfect Monitoring

Ayush Gupta*

25 October 2025

[Please click here for the latest version](#)

Abstract

This paper studies collusion between firms whose actions are constrained by behavioural or organizational concerns. I introduce a novel solution concept – pathwise ex post equilibrium (PXE) – for repeated games with imperfect public monitoring. In equilibrium, prescribed actions are required to be optimal for any possible realization of the signals. The resulting equilibria are ex post incentive compatible, have no punishment on the path of play, and are robust to different choice rules such as regret minimization and worst-case payoff maximization. When applied to a model with price setting firms, PXE yields sharp predictions about equilibrium outcomes: either firms play the same price in every period or prices follow a cyclical pattern consistent with the *rockets and feathers* price cycles documented in retail gasoline markets worldwide. I show that the observed cycles can serve as a robust and efficient mechanism for redistributing profits between firms and can therefore help sustain collusive agreements in the absence of explicit cash transfers. Price cycles may thus reflect deliberate choices rather than random shocks or coordination failures.

*Boston University, Department of Economics; guptaayu@bu.edu

I am indebted to Juan Ortner, Andrew Newman, Dilip Mookherjee, and Alex Chan for unbounded help and support throughout this project. I also thank Chengyang Zhu, Jakob von Ditzfurth, Andreas Haupt, and participants in the BU Theory seminars for many helpful comments and suggestions.

1 Introduction

A rich theoretical literature studies collusion between firms that compete over a long period of time. Existing results have yielded deep insights into cartel dynamics across a wide range of industries and institutional settings. However, most models in this literature assume that the firms are expected utility maximizers. In practice, colluding firms may have many other considerations. For example, agency problems may result in employees making decisions that don't maximize the firm's expected profit; ex post incentive compatibility may be necessary to foster trust and ensure the stability of a cartel; firms with external creditors may care about their worst-case profit because they wish to avoid bankruptcy. To study the impact of such considerations on collusive outcomes, it is necessary to depart from the standard framework and analyze firms that do not always maximize their expected profits.

This paper is motivated by a well documented phenomenon: retail gasoline prices in many cities exhibit a distinctive and persistent “rockets and feathers” pattern – a large and synchronous increase in pump prices is followed by many small and asynchronous decreases until prices approach marginal cost and the pattern repeats. Often called Edgeworth cycles or price cycles, this phenomenon has been studied extensively in the empirical literature but existing theories of competition and collusion struggle to offer a satisfactory explanation. Textbook Bertrand competition predicts marginal cost pricing, whereas collusion suggests a stable monopoly price. Competitive behaviour in the tradition of Maskin and Tirole (1988) can generate the observed price cycles, but only under strong assumptions which limit how often a firm can update its posted price. More general folk theorem constructions can accommodate almost any outcome, leaving unanswered the critical question of why this particular pattern emerges. These challenges suggest a gap in our understanding of collusion and call for a theoretical framework that can rationalize price cycles under assumptions consistent with the retail gasoline market.¹

In this paper, I model a repeated game with price setting firms and imperfect public monitoring.

¹A more detailed description of the observed price cycles and existing theories is contained in the next section.

I introduce a novel solution concept called *pathwise ex post equilibrium* (PXE). Roughly speaking, PXE requires that equilibrium actions are optimal for any possible realization of the public signal. Consequently, even in the worst-case outcome, equilibrium payoffs are higher than deviation payoffs. In the resulting equilibria, observed outcomes are constrained to one of two possibilities: either prices converge to a constant vector in finite time or they follow a cyclical pattern consistent with the price cycles observed in retail gasoline markets. I show that, within the set of PXE, price cycles can be efficient, serving as a compromise between two different equilibria with constant prices. Furthermore, when demand is sufficiently close to the Bertrand model of price competition, the optimal equilibrium for either firm is a price cycle.

My results are supported by Clark and Houde (2013) which documented evidence from the prosecution of a gas station cartel in Quebec, Canada. Former cartel members testified that they used asynchronous price changes to transfer surplus and ensure the stability of the cartel. That is precisely the mechanism identified in this paper. The firms use price changes to shift demand and thereby redistribute profits without explicit cash payments. The behavioural considerations formalized by PXE prevent the firms from raising their prices except when prices are sufficiently close to marginal cost. This forces most changes to be price cuts and results in the distinctive rockets and feathers pattern.

There are several motivations for a solution concept that explicitly incorporates non-expected utility considerations. First, due to imperfect monitoring, firms are often not certain if their rival has deviated or if they are observing an incorrect signal. So incentive compatibility requires the firms to sometimes initiate punishment even when no deviation has occurred. In practice, if the decision makers are employees, they may be very reluctant to start a costly price war when there is no conclusive evidence of cheating. As a result, punishments may be delayed, undermining the credibility of the collusive agreement.

Second, the possibility of a price war can itself undermine the viability of collusion. That is because even a brief price war can be extremely costly and firms with external creditors, risk-averse

managers, or short-term financial constraints may evaluate continuation payoffs using a worst-case criterion. In such cases, the risk of a price war may outweigh the expected gains from cooperation and make it impossible to sustain collusion, even when it is profitable on average.

Third, even in the absence of price wars, future payoffs are typically uncertain because they depend on the realization of noisy public signals. Studies such as Filiz-Ozbay and Ozbay (2007), Strack and Viefers (2021), Jhunjhunwala (2021), and Fioretti et al. (2022) suggest that human decision makers facing uncertainty don't necessarily maximize expected utility. They find evidence that decision makers often prioritize other objectives like minimizing future regret. This is particularly salient in my setting because collusion is illegal in most industries and therefore not enforceable via formal contracts. Sustaining cooperation under legal, reputational, and financial risks is thus likely to demand high levels of mutual trust and perceived fairness. In particular, a collusive agreement's ability to withstand future scrutiny may be just as important as its expected profitability. Colluding firms may therefore prefer strategies that are ex post incentive compatible, even if they do not maximize expected profits.

To study how such considerations shape equilibrium outcomes, I model firms that monitor their rivals using a randomly updated price comparison app. When the app updates, the observed price is equal to the rival's actual posted price. But when the app does not update, an outdated price is observed. Formally, at the end of period t , firm i observes a public signal y_j^t which equals the rival's actual price p_j^t with probability $1 - \gamma$ and repeats y_j^{t-1} with probability $\gamma \in (0, 1)$. The monitoring technology thus creates observational delays. If a firm changes its posted price, it may take a long time for the rival firm to learn the new price.

I assume that the firms evaluate continuation payoffs using the worst-case realization of future signals. After every public history, the firms play actions which minimize their worst-case loss from not deviating. The loss from not deviating is defined as the difference between a firm's realized profit and the profit it could have obtained by deviating. This reflects a preference for ex post optimality – the firms prefer actions that are optimal for any feasible realization of the public signals. The

solution concept formalizes this intuition. A pathwise ex post equilibrium (PXE) requires that there is no feasible realization of the public signals such that a firm could have benefited by deviating i.e. there is no loss from not deviating. That is why, even in the worst-case outcome, equilibrium payoffs are higher than the payoffs after any deviation.

Consequently, PXE are credible predictors of equilibrium outcomes under a variety of decision rules and heuristics that are used by humans in practice. In particular, firms that minimize regret and firms that maximize worst-case payoffs both find it optimal to not deviate from a strategy profile which is a PXE. Moreover, equilibrium actions are always ex post incentive compatible so even employees who are concerned about justifying their actions in the future would be willing to follow strategies that are a PXE.

Under the assumed monitoring technology, PXE requires that delaying a price change is never profitable. More specifically, if firm i is expected to change its price between periods t and $t + 1$, the new price cannot be less profitable than the old price. As a result, PXE do not require on-path punishments to ensure incentive compatibility. Firm i does not benefit by deviating and delaying the price change, so firm j does not need to provide incentives using continuation play. Even when the price comparison app does not update and the old price is observed, there is no need for firm j to initiate punishment.

In retail gasoline markets, the own price elasticity of demand facing an individual firm is typically very high. Consequently, price decreases are generally compatible with PXE but a price increase is feasible only when the old price is sufficiently close to marginal cost. Otherwise, the decrease in demand outweighs the increase in margin and makes a price increase unprofitable.

The behavioural considerations motivating PXE thus rationalize one key feature of price cycles – prices increase only when margins are sufficiently close to zero because earlier price increases are not ex post incentive compatible. That is because raising its price from an intermediate level is not myopically optimal for a firm. A collusive agreement must therefore use continuation play to ensure incentive compatibility. But continuation play is a function of the realized public signals

and hence uncertain. So the loss from not deviating is positive after a bad signal realization and the price increase will not remain optimal in hindsight.

I use this result to characterize the entire set of price paths that may be observed in a PXE. One possibility is that the vector of prices becomes constant in finite time. The only other possibility is that the prices exhibit a cyclical pattern consistent with the price cycles observed in retail gasoline markets. Formally, this cyclical pattern consists of countably many blocks of finite length. Within each block, prices are non-increasing and every block ends with a period of low prices. Crucially, all price increases occur at the transition between blocks, when prices are sufficiently low. So any sequence of prices which fits this pattern will closely resemble gasoline price cycles.

PXE thus rules out a large set of price paths. However, equilibria in which both firms play the same vector of prices in every period are not ruled out. Such constant price equilibria are a natural and intuitive benchmark when studying collusion between price setting firms. Any theory of price cycles must therefore explain why the firms do not coordinate on a constant cartel price. I posit that that is because the set of payoff vectors which can be achieved via constant prices is too limited. For example, under the Bertrand model of competition, the firms have equal market share in any equilibrium with constant prices. However, in practice, one of the firms may demand a larger share of the gains from cooperation.² Existing papers typically sidestep this issue by allowing for cash transfers or assuming that the firms can split demand in any ratio they want. However, neither assumption is reasonable in the retail gasoline market. Beyond their choice of posted price, gas stations have practically no control over where a consumer purchases gasoline. And since collusion is illegal in most markets, the firms have a strong incentive to avoid direct transfers if they carry an increased risk of detection and prosecution.

To formalize this intuition, I show that the firms can use a price cycle to achieve a vector of payoffs which cannot be achieved by any constant price equilibrium. By switching between periods with different relative prices, the firms shift demand and achieve an implicit transfer of profit, without

²A formal model of bargaining between the firms is beyond the scope of this paper. But it is intuitive that some firms may have more bargaining power by virtue of their size, geographic location, access to credit etc.

any explicit cash payments. The importance of such transfers for sustaining long term cooperation is already well documented in the literature. In the retail gasoline context, there is additional evidence, from a prosecuted cartel, supporting my claim that firms use price cycles precisely because they want to redistribute profits. Furthermore, Clark and Houde (2013) used data on prices and station characteristics to show that this mechanism can be used to transfer substantial amounts and can significantly reduce the frequency of deviations.

I also provide an explicit construction of strategies which are a PXE and generate the price cycles observed in retail gasoline markets. I show that there is no other equilibrium which can guarantee a higher payoff to both firms. The constructed price cycle is thus efficient and expands the set of equilibrium payoff vectors. Intuitively, the price cycle acts as a compromise between two distinct equilibria with constant prices. The price leader prefers the cycle over any equilibrium with equal prices because it is able to undercut the follower in some of the periods. The follower prefers the cycle over any equilibrium with unequal prices because it is able to match the leader in some of the periods.

To summarize, my results show that the firms use a price cycle when they wish to redistribute profits without explicit cash payments. The behavioural considerations formalized by PXE allow price increases only when the margins are sufficiently close to zero. This results in the distinctive rockets and feathers pattern observed in the data. Moreover, I show that price cycles are an efficient and robust transfer mechanism. They do not rely on the choice of parameters and they do not require assumptions about sticky prices or unobserved cost shocks.

This paper thus contributes to three distinct but related fields. First, I build on the rich empirical literature that has meticulously documented and analyzed price cycles across numerous markets all over the world. Foundational work in this area has established the key stylized facts that any successful theory must explain: the rockets and feathers pattern; adjustment delays during the undercutting phase; and the fact that prices increase only when they are close to marginal cost. While this body of work provides compelling evidence of the phenomenon, the underlying

mechanism remains a subject of active debate. My primary contribution here is to offer a novel explanation for the observed price cycles. I demonstrate why they might emerge as an equilibrium outcome in a setting with rational firms competing under a realistic form of imperfect monitoring.

Second, my work engages with the classic literature on collusion in a repeated game with imperfect monitoring. I study how the set of equilibrium outcomes changes when firms are guided by considerations beyond expected utility maximization. My focus is on behavioural and organizational concerns that shape decision making under uncertainty. But the equilibria I identify are robust and credible predictions of firm behaviour across a range of environments and selection criteria. The retail gasoline market provides a well-documented illustration of these mechanisms, but the framework applies more broadly to settings in which monitoring is imperfect, punishment is costly, and sustained collusion requires more than just positive expected payoffs.

Finally, I contribute to the literature on equilibrium refinements and behavioural game theory. The idea that human decision makers are motivated by a desire to avoid regret has a long history in economics. My solution concept – PXE – formalizes this concept within a repeated game with imperfect monitoring. In some ways, PXE is an extension of ex post perfect equilibria (XPE) as defined by Carroll (2024) and Krasikov and Lamba (2023). Whereas XPE is defined in a stochastic game with exogenous uncertainty, PXE is defined in a setting where endogenous actions affect the uncertainty. PXE provides a powerful selection criterion in any dynamic game where cooperation must be sustained under uncertainty and the players desire robust and stable agreements. Because equilibrium actions are always ex post optimal, PXE is particularly valuable for analyzing settings in which trust and perceived fairness are important. The solution concept also extends to games with imperfect private monitoring, making it a flexible tool for understanding cooperation in a wide range of environments.

2 Description of Price Cycles

Empirical data on gasoline price cycles dates back to at least 1993. Castanias and Johnson (1993) found evidence of price cycles in Los Angeles between 1968 and 1972. Since then, a plethora of papers have documented gasoline price cycles in many different cities at many different points in time. Some examples include Eckert (2003), Noel (2007a), and Noel (2007b) in Canada; Lewis (2012) in the United States; Foros and Steen (2013) in Norway; and Wang (2009a), Byrne (2012), de Roos and Katayama (2013), and Byrne and de Roos (2019) in Australia. A more complete survey can be found in Eckert (2013).

While there are some differences, the overall shape and structure of price cycles is remarkably consistent across time and space. Fig. 1 reproduces a graph from Byrne et al. (2025) to illustrate a typical gasoline price cycle.




Figure 1: Example of a Price Cycle

Price cycles are characterized by two distinct phases. In the undercutting phase, stations decrease their prices asynchronously and in small increments. Notably, some stations act as price leaders by

lowering their prices earlier and consequently enjoy elevated market shares for an extended period of time. This is puzzling because, in most jurisdictions, competing stations face minimal restrictions on how frequently they can update their own price. Stations can, and frequently do, adjust posted prices multiple times in a single day. So one would expect rapid price matching. Yet, in practice, competing stations often delay their response, allowing the price leaders to temporarily benefit from lower prices. This behaviour implies the existence of strategic considerations beyond simple short-run profit maximization.

In stark contrast, the restoration phase occurs abruptly and simultaneously. After a sustained period of incremental price cuts, retail prices collectively jump back to a higher common level. Crucially, this coordinated price restoration occurs only after margins approach marginal cost. Intermediate margin levels typically fail to trigger a price increase. This is also surprising because one expects an earlier price increase to raise profits for all firms.

In addition, it is important to keep in mind that gasoline is a largely homogeneous good. Individual gas stations face an extremely high own-price elasticity of demand – even small differences in prices often cause very large differences in market share. For example, Wang (2009b) estimates the elasticity to be as high as -18.77 for some gas stations. Given the high elasticity of demand, it is natural to expect that prices in a competitive market will be close to marginal cost. However, that is not the case in the retail gasoline market. As Fig. 1 illustrates, retail prices are frequently much higher than the wholesale price and changes in the two prices are often uncorrelated.

Maskin and Tirole (1988) showed that competitive behaviour can generate price cycles under Bertrand competition. However, their result relies on an alternating moves assumption under which each firm is only allowed to update its price in every other period. This assumption introduces a form of price stickiness that drives the slow and asynchronous nature of the undercutting phase. In periods when firm i is allowed to change its price, it undercuts firm j by the smallest increment possible. Firm i is able to do so because firm j is constrained to repeating its previous price and cannot respond to firm i 's anticipated price. In retail gasoline markets firms are typically allowed

to update their prices in every period and this mechanism breaks down. Firms respond to their rival's anticipated price and competition leads to marginal cost pricing.

Collusion based theories of price cycles also face significant challenges. In particular, if the firms are colluding, why are the prices frequently close to marginal cost? One would expect that the cartel can raise profits for both firms by effecting a price increase earlier in the cycle. On the surface, it appears suboptimal that prices increase only when margins are close to zero.

One possibility is that the firms are trying to sustain a high cartel price but imperfect monitoring leads to on-path price wars that manifest as price cycles. However, genuine price wars typically feature rapid and aggressive price cuts as firms actively attempt to undercut their rival. This is inconsistent with the observed price cycles which exhibit asynchronous price cuts and significant adjustment delays. On-path price wars in the spirit of Green and Porter (1984) thus do not appear to be a plausible explanation.

It is also worth noting that changes in marginal cost tend to be highly correlated in the retail gasoline industry. That is because the largest component of marginal cost is the wholesale price of gasoline, which is largely determined by global oil prices. Therefore, any theory which relies on unobserved cost shocks will struggle to model gasoline price cycles. In particular, a decrease in marginal cost may explain why a gas station lowers its price but it cannot explain why the rival delays matching the new price.

3 Model

I study a game with two firms indexed by $i \in \{1, 2\}$. Time is discrete and indexed by $t \in \{1, 2, \dots\}$. The firms have a common discount rate $0 < \delta < 1$. Both firms sell a single homogeneous good at a constant marginal cost given by c .

3.1 Actions and Signals

Define $p_i^t \in \mathcal{P}$ to be the actual price posted by firm i in period t . I assume that $\mathcal{P} = \{\epsilon, 2\epsilon, \dots, \bar{p}\}$ where $\epsilon > 0$ represents the smallest admissible change in price. \bar{p} can be interpreted as the consumers' maximum willingness to pay for the good. Assuming that prices are discrete is both realistic and analytically convenient. Retail prices are always quoted in finite increments and there are many practical concerns preventing infinitely small changes.³

Monitoring is imperfect so firms do not observe the actual price posted by their rival. Instead, at the end of each period, the firms observe a pair of independent public signals. The signal generating process is given by:

$$(y_1^1, y_2^1) = (p_1^1, p_2^1)$$

$$y_1^t = \begin{cases} p_1^t & \text{with probability } 1 - \gamma \\ y_1^{t-1} & \text{with probability } \gamma \end{cases}$$

$$y_2^t = \begin{cases} p_2^t & \text{with probability } 1 - \gamma \\ y_2^{t-1} & \text{with probability } \gamma \end{cases}$$

where $0 < \gamma < 1$ is fixed and common knowledge.⁴ I will use $y^t = (y_1^t, y_2^t) \in \mathcal{Y} \times \mathcal{Y}$ to denote the public signal profile observed at the end of period t . Note that $\mathcal{Y} = \mathcal{P}$.

This monitoring technology is consistent with a price comparison app that updates at random times. When a firm observes its rival's price on the app, it does not know if the app updated in that period. So the observed price may be accurate or it may be an older price. In particular, when a firm changes its posted price, the rival does not always learn the new price immediately. This informational friction captures a common challenge faced by firms in many industries – changes in actions are often observed with a delay.

³In the United States, ϵ is typically one-tenth of a cent for gas stations.

⁴One signal may update while the other does not. So the probability that both signals are accurate is $(1 - \gamma)^2$.

3.2 Demand

In order to focus on collusion between the firms, I abstract away from a formal model of consumer choice. However, I assume that the demand for firm i in period t is a function of p_i^t , y_i^t , and y_j^t .

One interpretation is that consumers don't observe the posted prices so they use the public signals to decide where to shop. Demand is thus a function of both public signals. If they make a purchase, consumers have to pay the actual price posted by their chosen firm. Consumers learn the posted price before completing their purchase and can alter their consumption decision if the actual price is different from the observed signal. Demand for firm i is therefore affected by p_i^t as well.

Formally, the demand for firm i in period t is given by $D(p_i^t, y_i^t, y_j^t)$ and I assume:

Assumption 1. $D_i(p_i^t, y_i^t, y_j^t)$ is uniformly bounded by M .

Assumption 2. $D_i(p_i^t, y_i^t, y_j^t)$ is decreasing in y_i^t and increasing in y_j^t .

Assumption 3. There exists $\eta > 0$ such that when $p_i^t \neq y_i^t$:

$$D_i(p_i^t, y_i^t, y_j^t) \geq (1 + \eta)D_i(y_i^t, y_i^t, y_j^t) \text{ if } p_i^t < y_i^t.$$

$$D_i(p_i^t, y_i^t, y_j^t) \leq (1 - \eta)D_i(y_i^t, y_i^t, y_j^t) \text{ if } p_i^t > y_i^t.$$

Assumption 2 states that the demand for firm i is decreasing in its own signal and increasing in its rival's signal. The implicit assumption is that consumers prefer to shop at the lower priced firm but the firms are spatially differentiated. So some consumers will continue to shop at firm i even when its observed price is higher.

Assumption 3 states that the demand for firm i is affected by p_i^t even when the public signals don't update. That is because consumers pay p_i^t and can alter their consumption decision if the posted price turns out to be different from the observed signal.

The stage payoff for firm i in period t is given by:

$$u_i(p_i^t, y_i^t, y_j^t) = D(p_i^t, y_i^t, y_j^t)(p_i^t - c)$$

Note that the demand and stage payoff functions are the same for both firms. However, I use subscripts to clarify which firm's demand / payoff is being referred to.

Remark 1. Given the monitoring technology and the stage payoff function, this framework can also be interpreted as a stochastic game with an action dependent state variable. Under that interpretation, the state variable in period t is the vector of public signals observed at the end of period $t - 1$. The vector of public signals determines the mapping from actions to payoffs and is itself a function of the actions played in the previous period.

3.3 Histories and Strategies

Define $h^t = (y^1, y^2, \dots, y^{t-1})$ to be the public history at the beginning of period t . Let \mathcal{H}^t be the set of all public histories of length $t - 1$. Then \mathcal{H}^∞ is the set of all infinitely long public histories with generic element $h = (y^1, y^2, \dots)$.

Let s_i be a public strategy for firm i . Then s_i is a measurable mapping from the set of public histories to the set of actions. Throughout this paper, I restrict attention to public strategies.

Let $\mathcal{Y}_i(s_i | h^t)$ be the set of firm i signals that are consistent with the public strategy s_i and the public history h^t . That means every element of $\mathcal{Y}_i(s_i | h^t)$ is in the support of y_i^t when the firm i plays according to s_i after h^t is realized.

I say that $h \in \mathcal{H}^\infty$ is *consistent* with the public strategy profile $s = (s_1, s_2)$ if, for every t , there is a positive probability that the public history at the beginning of period t is given by the first $t - 1$ elements of h when the firms play according to s .

I use $\mathcal{H}^\infty(s)$ to denote the set of all infinitely long histories that are consistent with s . I say that $h \in \mathcal{H}^\infty(s \mid h^t)$ if for every $\tau > t$, there is a positive probability that the public history at the beginning of period τ is given by the first $\tau - 1$ elements of h when the firms play according to the continuation strategy profile $s \mid h^t$ after h^t is realized.

3.4 Payoffs and Preferences

Fix any public history h^t . Assume that the firms play according to the public strategy profile s . Then the *guaranteed* payoff for firm i at the beginning of period t is given by:

$$g_i^t(s \mid h^t) = \inf_{h \in \mathcal{H}^\infty(s \mid h^t)} (1 - \delta) \sum_{\tau \geq t} \delta^{\tau-t} u_i(s_i(h^\tau), y_i^\tau, y_j^\tau)$$

where (y_i^τ, y_j^τ) is the τ element of $h \in \mathcal{H}^\infty(s \mid h^t)$ and h^τ is equal to the first $\tau - 1$ elements of h .

Lemma 1. The following exists for every i , every s , and every h^t :

$$\min_{h \in \mathcal{H}^\infty(s \mid h^t)} (1 - \delta) \sum_{\tau \geq t} \delta^{\tau-t} u(s_i(h^\tau), y_i^\tau, y_j^\tau)$$

Since the minimum always exists, I can redefine guaranteed payoffs using the minimum instead of the infimum. The proof is contained in the appendix but the intuition is that the set of infinitely long histories is compact. The subset that is consistent with s is also compact and the objective function is continuous over that subset. And the minimum of a continuous function on a compact set always exists.

Notice that the guaranteed payoffs at time t depend only on the public history and the choice of continuation strategies. Moreover, there is no ambiguity so each firm's prior is a singleton and the rectangularity condition from Epstein and Schneider (2003) is trivially satisfied. As a result, guaranteed payoffs can be written as:

$$g_i^t(s \mid h^t) = \min_{(y_i^t, y_j^t) \in \mathcal{Y}(s \mid h^t)} \left\{ (1 - \delta) u(s_i(h^t), y_i^t, y_j^t) + \delta g_i^{t+1}(s \mid h^t, y_i^t, y_j^t) \right\} \quad (1)$$

That is because:

$$g_i^{t+1}(s \mid h^t, y_i^t, y_j^t) = \min_{h \in \mathcal{H}^\infty(s \mid h^t, y_i^t, y_j^t)} (1 - \delta) \sum_{\tau > t} \delta^{\tau-t} u(s_i(h^\tau), y_i^\tau, y_j^\tau)$$

Notice that $h \in \mathcal{H}^\infty(s \mid h^t)$ if $h \in \mathcal{H}^\infty(s \mid h^t, y_i^t, y_j^t)$ and $(y_i^t, y_j^t) \in \mathcal{Y}(s \mid h^t)$. So the RHS of (1) can be written as:

$$\begin{aligned} &= \min_{(y_i^t, y_j^t) \in \mathcal{Y}(s \mid h^t)} \left\{ (1 - \delta) u(s_i(h^t), y_i^t, y_j^t) + \min_{h \in \mathcal{H}^\infty(s \mid h^t, y_i^t, y_j^t)} (1 - \delta) \sum_{\tau > t} \delta^{\tau-t} u(s_i(h^\tau), y_i^\tau, y_j^\tau) \right\} \\ &= \min_{h \in \mathcal{H}^\infty(s \mid h^t)} (1 - \delta) \sum_{\tau \geq t} \delta^{\tau-t} u(s_i(h^\tau), y_i^\tau, y_j^\tau) \\ &= g_i^t(s \mid h^t) \end{aligned}$$

Remark 2. Dynamic consistency is satisfied because, conditional on s and h^t , the value of $g_i^{t+1}(s \mid h^t)$ is a deterministic function of (y_i^t, y_j^t) . So the guaranteed payoffs computed at the beginning of period t are equal to the guaranteed payoffs computed at the beginning of period $t + 1$. In particular, information revealed in period t does not change the value of any $g_i^{t+1}(s \mid h^t, y_i^t, y_j^t)$.

With a slight abuse of notation I define:

$$g_i^t(s \mid h^t, y_j^t) = \min_{y_i^t \in \mathcal{Y}_i(s_i \mid h^t)} \left\{ (1 - \delta) u_i(s_i(h^t), y_i^t, y_j^t) + \delta g_i^{t+1}(s \mid h^t, y_i^t, y_j^t) \right\}$$

Now define:

$$r_i(s, h^t) = \max_{y_j^t \in \mathcal{Y}_j(s_j \mid h^t)} \max_{\hat{s}_i \neq s_i} \left\{ g_i^t(\hat{s}_i, s_j \mid h^t, y_j^t) - g_i^t(s_i, s_j \mid h^t, y_j^t) \right\}$$

Then $r_i(s, h^t)$ can be interpreted as the maximum *regret* experienced by firm i at the end of period t when both firms play according to the public strategy profile s . It is the maximum payoff that firm i can gain by deviating in period t . I assume that the firms want to minimize $r_i(s, h^t)$. So

after h^t is realized, the optimization problem for firm i is to pick a strategy s_i such that:

$$s_i \in \arg \min_{s'_i} r_i(s'_i, s_j, h^t)$$

Remark 3. When the firms are expected utility maximizers, they minimize expected regret instead of maximum regret.

Remark 4. The realization of y_j^t depends only on the public history h^t and firm j 's strategy s_j . Neither is affected by firm i 's action in period t . That is why firm i treats the realization of y_j^t as fixed when evaluating regret at the end of period t .

3.5 Equilibria

Definition 1. A profile of public strategies s is a *pathwise ex post equilibrium* (PXE) if for both i and every public history h^t :

$$r_i(s, h^t) \leq 0$$

In practice, this definition can be difficult to work with. So I will prove that the following definition is equivalent and use it for most of my subsequent results:

Definition 2. A profile of public strategies s is a PXE if for every i , every public history h^t , and every $y_j^t \in \mathcal{Y}_j(s_j \mid h^t)$:

$$\min_{y_i^t \in \mathcal{Y}_i(s_i \mid h^t)} \left\{ (1 - \delta) u_i(a_i^t, y_i^t, y_j^t) + \delta g_i^{t+1}(s \mid h^t, y_i^t, y_j^t) \right\} \geq \min_{\hat{y}_i^t \in \mathcal{Y}_i(\hat{s}_i \mid h^t)} \left\{ (1 - \delta) u_i(\hat{a}_i^t, \hat{y}_i^t, y_j^t) + \delta g_i^{t+1}(s \mid h^t, \hat{y}_i^t, y_j^t) \right\} \quad (2)$$

where $a_i^t = s_i(h^t)$ is the prescribed action for firm i and $\hat{a}_i^t = \hat{s}_i(h^t)$ is a one shot deviation.

It is obvious that Definition 1 implies Definition 2 so I only prove the converse:

Lemma 2. A public strategy profile s is a PXE if it satisfies the condition in Definition 2.

The left-hand side of Condition (2) is firm i 's guaranteed payoff in equilibrium. The right-hand side is its guaranteed payoff after a deviation in period t . Both terms condition on the same realization of y_j^t . Thus, PXE requires that, for any realization of y_j^t , the guaranteed payoff in equilibrium is no less than the guaranteed payoff after a deviation.

PXE is closely related to ex post perfect equilibrium (XPE) as defined by Carroll (2024) and Krasikov and Lamba (2023). Like XPE, PXE requires that equilibrium actions are ex post incentive compatible. In other words, the prescribed actions in period t must be optimal for any possible realization of the public signals in period t . Intuitively, ex post incentive compatibility means that actions are dominant with respect to the set of public signals that may be realized in equilibrium. This requirement is much stronger than standard sequential rationality, which only requires that actions are incentive compatible in expectation.

However, PXE and XPE differ in some subtle but important ways. XPE is defined in a stochastic game framework, where the uncertainty is over the exogenous realization of future stage games. In contrast, the uncertainty in my setting is over the realization of future public signals, which depend on the firms' endogenous actions.

Because the realization of public signals depends on endogenous choices, formalizing ex post incentive compatibility is challenging in this setting. In equilibrium, firms never deviate, so the sequence of public signals following a deviation is never observed. Incentive compatibility therefore relies on beliefs about what would have happened if a deviation had occurred. PXE assumes that firms evaluate deviation payoffs under the worst possible realization of public signals. This is ultimately a modelling choice and an argument can be made for alternative criteria. However, most subsequent results are unchanged if firms instead use the *expected* sequence of public signals.

PXE's restrictions make it a credible predictor of behaviour in many settings. Because equilibrium actions are ex post incentive compatible, PXE are easy to sustain even when the decision maker is an employee who might have to justify her choices in the future. She can be certain that the prescribed action will appear optimal, regardless of which public signals are realized. So the employee knows that she won't be penalized for following the equilibrium strategy.

A related benefit is that firms playing a PXE experience no regret on the path of play. Because equilibrium actions are optimal for any possible realization of the public signals, there is no realization such that a deviation would have yielded a higher payoff. Consequently, the firms never lose payoff by not deviating and never wish to revise past actions.

Finally, PXE ensures that the worst-case payoffs in equilibrium exceed those after any deviation. If the equilibrium requires a firm to sacrifice some payoff today, it knows that, even in the worst case, future compensation will offset today's loss. Consequently, even firms that maximize worst-case payoffs, instead of expected payoffs, will not deviate if the strategy profile is a PXE.

Remark 5. Any dominant strategy equilibrium of the stage game is always a PXE. For example, when the stage game is prisoner's dilemma, it is a PXE for both players to defect after every history.

In general, a mixed strategy equilibrium of the stage game cannot be supported in a PXE. For example, when the stage game is rock paper scissors, there are no off-punishments that make it ex post incentive compatible for both players to randomize over their three actions. Intuitively, if player 1's randomization results in rock and player 2's randomization results in paper, player 1 regrets not deviating and playing scissors.

Definition 3. An equilibrium with strategy profile s is *dominated* by an equilibrium with strategy profile s' if $g_i(s') > g_i(s)$ and $g_j(s') \geq g_j(s)$ for some i .

An equilibrium is *efficient* if it is not dominated by any other equilibrium.

Definition 4. An equilibrium is *stationary* if $s(h^t) = p = (p_i, p_j)$ for all on-path histories.

Note that a stationary equilibrium need not have $p_i = p_j$.

4 General Results

All proofs are included in the appendix but I provide some intuition where appropriate.

Lemma 3. Fix a mixed strategy equilibrium σ . There exists a pure strategy equilibrium s such that $g(s) = g(\sigma)$.

I will restrict attention to pure strategies for the remainder of this paper. Lemma 3 establishes that doing so is without loss of generality. The intuition for Lemma 3 comes from the fact that the guaranteed payoffs in any equilibrium are determined by the worst-case sequence of public signals. If σ is a mixed strategy equilibrium, the worst-case sequence of signals under σ must be consistent with some pure strategy profile that is in the support of σ . The proof of Lemma 3 shows that this pure strategy profile can also be sustained in equilibrium and can therefore replicate the guaranteed payoffs under σ .

Proposition 1. Fix a strategy profile s and a public history h^t . Then s is a PXE only if:

$$s_i(h^t) = p_i^t \neq y_i^{t-1} \implies u_i(p_i^t, y_i^{t-1}, y_j^t) \geq u_i(y_i^{t-1}, y_i^{t-1}, y_j^t)$$

where y_j^t is any public signal which is realized with positive probability.

Proposition 1 formalizes the main restriction imposed by PXE on equilibrium strategies. A firm will not change its price if the new price is less profitable than the old price. Conversely, if the equilibrium requires a firm to update its price, the new price must yield a higher stage payoff than the old price.

Suppose this condition is not satisfied – the equilibrium requires firm i to update its price but p_i^t yields a strictly lower stage payoff. Firm i has an incentive to deviate and delay the price change i.e. play y_i^{t-1} again. To play p_i^t and sacrifice payoff in period t , firm i must be incentivized using continuation play. However, any sequence of public signals which may be realized after y_i^{t-1} may also be realized after firm i plays p_i^t . That is because the price comparison app can fail to update and display y_i^{t-1} even when firm i plays p_i^t . So the worst-case continuation payoff after playing p_i^t cannot be higher than the worst-case continuation payoff after deviating to y_i^{t-1} . This violates ex post incentive compatibility since the total guaranteed payoff for firm i is strictly higher if it deviates.

Remark 6. In equilibrium, p_i^t and y_i^{t-1} are the only signals that are observed with positive probability. Proposition 1 establishes that p_i^t always yields weakly higher stage payoff. So firm i already has an incentive to play p_i^t over y_i^{t-1} . Firm j does not need to provide incentives using continuation play. In particular, firm i will not deviate even if its continuation payoffs are the same under both signal realizations. That is why I say that there is no punishment on the path of play in a PXE.

Lemma 4. There exists $\underline{p} \in P$ such that if $y_i^{t-1} > \underline{p}$ then $y_i^{t-1} \geq p_i^t$ in every PXE.

When demand is elastic, it is profitable for a firm to raise its price only when the old price is sufficiently close to marginal cost. Otherwise, the decrease in demand outweighs the increase in margin and profit decreases. Lemma 4 formalizes this intuition and derives an upper bound on the old price. When its old price is greater than \underline{p} , it is never profitable for a firm to raise its price. So the new price must be weakly lower. I stress that the derived bound is a function of the elasticity parameter η . When η is close to zero, the restriction imposed by Lemma 4 may be trivial. But when demand is very elastic, as is typically the case in the retail gasoline market, \underline{p} will be close to marginal cost.⁵

Lemma 4 thus rationalizes one of the key features of the observed price cycles – prices increase


⁵I remind the reader that some studies estimate the own price elasticity of demand to be as high as -18.77 for gas stations.

only when they are sufficiently close to marginal cost. The reason being that the behavioural considerations formalized by PXE rule out price increases from intermediate levels. So while they may be inefficient in an unconstrained setting, the observed strategies can be optimal given the PXE constraints.

Example 1. Consider the following profile of on-path strategies:

$$p_1^t = \begin{cases} 10 & \text{if } t \text{ is odd} \\ 9 & \text{if } t \text{ is even} \end{cases} \quad \text{and} \quad p_2^t = 10$$

When the firms play according to these strategies, one possible sequence of observed prices is:



When $c = 0$ and η is sufficiently large, the on-path strategies in Example 1 *cannot* be supported in a PXE. That is because $p_i^t = 9$ is more profitable than $p_i^t = 10$ for firm 1. So there are no off-path strategies that can make it ex post incentive compatible for firm 1 to raise its price from 9 to 10.

The significance of Example 1 is that it shows how PXE can rule out an intuitive way for the firms to redistribute profits. When the prices are already high, PXE prevents further increases. So any change in market share requires at least one firm to lower its price. Consequently, when the prices are high, the firms have only two choices. Either they can maintain the current price and get the same vector of payoffs in every subsequent period; or at least one firm must lower its price to a level which might appear suboptimal.

Remark 7. In a more general setting, PXE restricts transitions between action profiles. In particular, if it is feasible to switch from an action profile a to a different action profile a' , it is typically not feasible to switch from a' to a . That is because of the restriction imposed by Proposition 1. If a player is willing to switch from a to a' then she must prefer a' over a . But that means she does not prefer a over a' and is therefore unwilling to make that switch.⁶

Proposition 2. Fix any $h \in \mathcal{H}^\infty(s)$ where s is a PXE.

Then either there exists some T such that $y^t = y^T$ for every $t \geq T$.

Or there exists a strictly increasing sequence of natural numbers $\{\tau_n\}$ such that:

1. $\tau_1 = 1$.
2. y^t is non-increasing when $\tau_n \leq t < \tau_{n+1}$.
3. There exists i such that $y_i^t \leq \underline{p}$ when $t = \tau_n - 1$ for any $n \geq 2$.

Note that y^t is a vector. So y^t non-increasing means that $y_i^{t+1} \leq y_i^t$ for all i .

Proposition 2 uses Lemma 4 to characterize the set of price paths which may be observed in a PXE. It establishes that any sequence of observed prices will be (eventually) stationary or it will follow the cyclical pattern observed in gasoline markets. More specifically, any non-stationary sequence of prices can be split into a countable number of blocks such that, within each block, prices are non-increasing and all prices increases happen at the transition between blocks, when prices are sufficiently low.

⁶The only exception is when the player changing her action is indifferent between the two action profiles.

5 Efficiency of Price Cycles

The results in Section 4 provide some intuition for why price cycles are observed in retail gasoline markets. By changing their posted price, the firms are able to shift demand and change the distribution of profits. But the constraints imposed by PXE do not allow for arbitrary changes in price. In particular, prices can increase only when they are sufficiently close to marginal cost. This forces most changes to be price cuts and leads to the distinctive rockets and feathers pattern. However, equilibria with a stationary cartel price are not ruled out by PXE. So this section shows why colluding firms may prefer a price cycle over a stationary price.

To fix ideas, consider the following example:

Example 2. Let on-path strategies be given by:

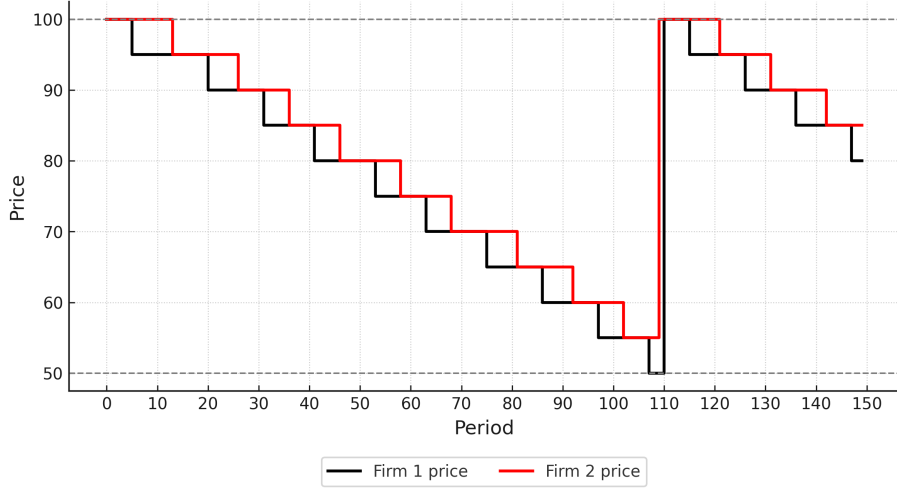
$$\begin{aligned}
 (p_1^1, p_2^1) &= (\bar{p}, \bar{p}) \\
 p_1^t &= \begin{cases} \bar{p} & \text{if } (y_1^{t-1}, y_2^{t-1}) = (p^{min}, p^{min}) \\ y_1^{t-1} - \epsilon & \text{if } y_1^\tau = y_2^\tau \text{ for all } t - k_1 \leq \tau \leq t - 1 \\ y_1^{t-1} & \text{otherwise} \end{cases} \\
 p_2^t &= \begin{cases} \bar{p} & \text{if } (y_1^{t-1}, y_2^{t-1}) = (p^{min}, p^{min}) \\ y_1^{t-1} & \text{if } p^{min} < y_1^\tau < y_2^\tau \text{ for all } t - k_2 \leq \tau \leq t - 1 \\ y_2^{t-1} & \text{otherwise} \end{cases}
 \end{aligned}$$

Fig. 2 shows a sequence of prices that may be observed when the firms play according to the strategies in Example 2.

In this example, both firms start by posting \bar{p} as their price in the first period.⁷ The firms post the same price in the first k_1 periods of the game. In period $k_1 + 1$, firm 1 undercuts its rival and lowers its price to $\bar{p} - \epsilon$. Firm 2 does not react immediately and continues to post \bar{p} . After k_2 periods of

⁷Recall that \bar{p} is the upper bound on the set of prices.

Figure 2: Sequence of Observed Prices



unequal prices, firm 2 also lowers its price by ϵ and both firms post $\bar{p} - \epsilon$. This process repeats until both firms start posting p^{min} . At that point, both firms raise their price back up to \bar{p} .

As is clear from Fig. 2, the resulting sequence of prices matches the price cycles observed in retail gasoline markets. So I will start by proving that the on-path strategies in Example 2 can be supported in a PXE when the following assumptions are satisfied:

Assumption 4. $D_i(p_i^t, y_i^t, y_j^t) > 0$ for every p_i^t .

Then there exists $\psi \in (\eta, 1)$ such that for every $p_i^t > y_i^t$:

$$D_i(p_i^t, y_i^t, y_j^t) \geq (1 - \psi)D_i(y_i^t, y_i^t, y_j^t)$$

Assumption 5. $\epsilon < \eta(1 - \psi)(\bar{p} - c)$.

First notice the contrast between Assumption 3 and Assumption 4. Assumption 3 requires a lower bound on the decrease in demand when firm i raises its price. Assumption 4 requires an upper bound. Collectively, the two assumptions require that any increase in price results in strictly lower

demand but demand never falls to zero.

Assumption 5 requires that ϵ is sufficiently small i.e. the grid of possible prices is not too sparse.

Given these assumptions, I can pick $p^{min} > c$ such that:

$$\begin{aligned} p^{min} &\leq c + (1 - \psi)(\bar{p} - c) \\ p^{min} &\geq c + \frac{\epsilon}{\eta} \end{aligned}$$

Lemma 5. The on-path strategies in Example 2 can be supported in a PXE whenever δ is sufficiently close to 1.

The formal proof is contained in the appendix but the off-path punishments are very simple. If firm i deviates and posts a price that is lower than the prescribed price, firm j responds by lowering its own price. Firm i 's demand is decreasing in y_j^t so the threat of punishment lowers firm i 's guaranteed payoff and deters deviations.

The equilibrium in Example 2 is also efficient under some mild assumptions. First define:

$$\begin{aligned} D(y_1^t, y_2^t) &= D_1(y_1^t, y_1^t, y_2^t) + D_2(y_2^t, y_1^t, y_2^t) \\ \sigma_i(y_1^t, y_2^t) &= \frac{D_i(y_i^t, y_1^t, y_2^t)}{D(y_1^t, y_2^t)} \end{aligned}$$

So $D(y_1^t, y_2^t)$ is the total market demand when the signals are accurate and $\sigma_i(y_1^t, y_2^t)$ is the market share for firm i .

Assumption 6. $D(y_1^t, y_2^t)$ is strictly concave on $\mathcal{Y} \times \mathcal{Y}$.

Assumption 7. The slope of $\sigma_i(y_1^t, y_2^t)$ with respect to y_i^t is weakly decreasing in $|y_1^t - y_2^t|$.

Assumption 6 is a standard concavity assumption. It ensures the existence of a unique profit maximizing monopoly price. Assumption 7 requires that market shares are most sensitive to price

changes when both firms are posting the same price. It is consistent with a world in which some consumers are more price sensitive than others. When firm i undercuts its rival by ϵ , a mass of consumers shift their demand from firm j to firm i . The consumers who continue to purchase from firm j are less price sensitive so further price cuts by firm i result in smaller changes in market shares. Intuitively, Assumption 7 is satisfied when the demand for firm i , conditional on the price of firm j , has a sigmoid shape centred on firm j 's price.

Notice that Assumption 7 is a weaker version of the assumption imposed by the Bertrand model of price competition. Under Bertrand competition, the firm with the lower price serves the entire market. So σ_i is vertical when $y_1^t = y_2^t$ and horizontal everywhere else. Such an extreme form of demand is consistent with Assumption 7 but is unnecessary for my results.

I stress that Assumption 6 and Assumption 7 are needed only to ensure that the specific equilibrium in Example 2 is efficient. All other results are valid when these assumptions do not hold. Moreover, even when Assumption 6 or Assumption 7 is violated, there typically exists a price cycle equilibrium which is efficient. In that scenario, the key difference will be that the firms may lower their price by more than ϵ at a time.

Proposition 3. The equilibrium in Example 2 is efficient.

Intuitively, a price cycle can be efficient because it allows the firms to switch between periods with different relative prices. If, instead, the firms play a stationary equilibrium with equal prices, firm 1's payoff will be strictly lower because it will not enjoy any periods with extra demand. Similarly, if the firms play a stationary equilibrium with different prices, firm 2's payoff will be strictly lower because there will be no periods with equal demand.

The proof of Proposition 3 relies on the fact that PXE rules out arbitrary changes in price. In particular, it is not ex post incentive compatible for firm 1 to raise its price when the public signals are $(\bar{p} - \epsilon, \bar{p})$. That is why, in the constructed example, the sequence of prices is:

$$(\bar{p}, \bar{p}); (\bar{p} - \epsilon, \bar{p}); (\bar{p} - \epsilon, \bar{p} - \epsilon); \dots$$

and not:

$$(\bar{p}, \bar{p}); (\bar{p} - \epsilon, \bar{p}); (\bar{p}, \bar{p}); \dots$$

The alternate sequence is more profitable for both firms but it is not within the set of PXE outcomes. That is why it does not dominate the constructed example.

The next step is proving that it is efficient for firm 1 to lower its price in the first place. When firm 2 eventually matches the price cut, both firms end up with the same market shares but at a lower price. Nevertheless, it is profitable for firm 1 to lower its price if the phase with unequal prices is sustained for sufficiently long. The reason is that future payoffs are discounted. So the extra demand today is more valuable to firm 1 than the lost margin in the future. The same argument can be used to show why firm 2 matches firm 1's price, even though that causes firm 1 to eventually lower its price further.

Perhaps the most striking feature of this result is the fact that the constructed equilibrium is efficient even though the firms play $(\bar{p} - \epsilon, \bar{p} - \epsilon)$, which is strictly dominated (in the static sense) by (\bar{p}, \bar{p}) . The reason is that any additional periods with the prices at (\bar{p}, \bar{p}) must come at the beginning of the game because of the irreversibility imposed by PXE. Doing so delays the first period in which firm 1 can undercut its rival. Because of discounting, this delay causes a discrete drop in firm 1's payoff. By instead switching from $(\bar{p} - \epsilon, \bar{p})$ to $(\bar{p} - \epsilon, \bar{p} - \epsilon)$ later in the game, the firms can achieve a smaller transfer of payoff from firm 1 to firm 2.

In effect, the set of payoffs that can be guaranteed with only two distinct price vectors is discrete and sparse. Playing a third price vector later in the game allows the firms to make smaller adjustments to their relative payoffs. The equilibrium therefore trades a small amount of total surplus for finer control over the distribution of surplus across firms. Under PXE, this trade-off preserves overall efficiency. No alternative sequence of prices can make both firms better off.

6 Conclusion

This paper introduces a novel solution concept – pathwise ex post equilibrium (PXE) – to study how behavioural and organizational concerns shape equilibrium outcomes in a repeated game with imperfect monitoring. My results suggest that the rockets and feathers price cycles commonly observed in retail gasoline markets may not be caused by competitive undercutting or random shocks. Instead, they can be a deliberate choice, serving as a robust and efficient mechanism for the firms to redistribute profits and sustain collusive agreements.

The model assumes a realistic form of imperfect monitoring that reflects the realities of the retail gasoline market – firms observe their rival’s price through randomly updated and potentially delayed public signals, such as online price-tracking applications. Under these conditions, PXE imposes a critical restriction on firm behaviour: a firm will only change its price if the new price is guaranteed to be at least as profitable as the old one, regardless of the signal observed by its rival.

This restriction generates the key dynamics of a price cycle. When demand is elastic, price decreases are generally compatible with PXE but price increases are feasible only when the current price is sufficiently close to marginal cost. In particular, small price increases from intermediate levels are ruled out. This finding explains both the sharp, coordinated price restorations from a low base (the rockets) and the subsequent phase of slow, asynchronous price decreases (the feathers). Any pathwise ex post equilibrium must therefore feature constant prices or exhibit this cyclical pattern.

Crucially, my results establish that price cycles are not just feasible but can also be Pareto efficient. If the firms differ in their bargaining power, a stationary cartel price may not be agreeable to all parties. Price cycles provide a dynamic mechanism to redistribute profits over time without resorting to illegal cash payments. By allowing for periods of price leadership where one firm enjoys a temporary market share advantage, the cycle facilitates a division of collusive profits that can be more desirable for all firms. This theoretical conclusion is supported by empirical evidence from the prosecution of a gas station cartel – former cartel members testified that adjustment delays were

used specifically to redistribute surplus within the cartel.

In conclusion, this paper makes a significant contribution by demonstrating that the puzzling and persistent phenomenon of gasoline price cycles can be understood as a rational, robust, and efficient collusive strategy. By integrating a behaviourally motivated solution concept with a realistic form of imperfect monitoring, the analysis furthers our understanding of how non-expected utility considerations affect collusive outcomes in an uncertain environment. The PXE solution concept is itself a valuable theoretical tool with broader applicability. It can be used for analyzing cooperation in any setting where the players may be concerned about ex post incentive compatibility, regret minimization, or worst-case payoffs.

References

- David Blackwell. Discounted dynamic programming. *Annals of Mathematical Statistics*, 36(1): 226–235, 1965.
- David Byrne. Petrol price cycles. *Australian Economic Review*, 45(4):497–506, 2012.
- David P. Byrne and Nicolas de Roos. Learning to coordinate: A study in retail gasoline. *American Economic Review*, 109(2):591—619, 2019.
- David P. Byrne, Nicolas de Roos, Matthew S. Lewis, Leslie M. Marx, and Xiaosong Wu. Price coordination with asymmetric information sharing: Theory and evidence. Working Paper, 2025.
- Gabriel Carroll. Dynamic incentives in incompletely specified environments. Working Paper, 2024.
- Rick Castanias and Herb Johnson. Gas wars: Retail gasoline price fluctuations. *The Review of Economics and Statistics*, 75(1):171–174, 1993.
- Robert Clark and Jean-François Houde. Collusion with asymmetric retailers: Evidence from a gasoline price-fixing case. *American Economic Journal: Microeconomics*, 5(3):97–123, 2013.

- Nicolas de Roos and Hajime Katayama. Gasoline price cycles under discrete time pricing. *The Economic Record*, 89(285):175–193, 2013.
- Andrew Eckert. Retail price cycles and the presence of small firms. *International Journal of Industrial Organization*, 21(2):151–170, 2003.
- Andrew Eckert. Empirical studies of gasoline retailing: A guide to the literature. *Journal of Economic Surveys*, 27(1):140–166, 2013.
- Larry G. Epstein and Martin Schneider. Recursive multiple-priors. *Journal of Economic Theory*, 113(1):1–31, 2003.
- Emel Filiz-Ozbay and Erkut Y. Ozbay. Auctions with anticipated regret: Theory and experiment. *American Economic Review*, 97(4):1407–1418, 2007.
- Michele Fioretti, Alexander Vostroknutov, and Giorgio Coricelli. Dynamic regret avoidance. *American Economic Journal: Microeconomics*, 14(1):70–93, 2022.
- Oystein Foros and Frode Steen. Vertical control and price cycles in gasoline retailing. *The Scandinavian Journal of Economics*, 115(3):640–661, 2013.
- Edward J. Green and Robert H. Porter. Noncooperative collusion under imperfect price information. *Econometrica*, 52(1):87–100, 1984.
- Tanushree Jhunjhunwala. Searching to avoid regret: An experimental evidence. *Journal of Economic Behavior & Organization*, 189:298–319, 2021.
- Ilia Krasikov and Rohit Lamba. Uncertain repeated games. Working Paper, 2023.
- Matthew S. Lewis. Price leadership and coordination in retail gasoline. *Journal of Urban Economics*, 72(3):320–334, 2012.
- Eric Maskin and Jean Tirole. A theory of dynamic oligopoly ii: Price competition, kinked demand curves, and edgeworth cycles. *Econometrica*, 56(3):571–599, 1988.

- Michael D. Noel. Edgeworth price cycles: Evidence from the toronto retail gasoline market. *The Journal of Industrial Economics*, 55(1):69–92, 2007a.
- Michael D. Noel. Edgeworth price cycles, cost-based pricing, and sticky pricing in retail gasoline markets. *The Review of Economics and Statistics*, 89(2):324–334, 2007b.
- Philipp Strack and Paul Viefers. Too proud to stop: Regret in dynamic decisions. *Journal of the European Economic Association*, 19(1):165–199, 2021.
- Zhongmin Wang. (mixed) strategy in oligopoly pricing: Evidence from gasoline price cycles before and under a timing regulation. *Journal of Political Economy*, 117(6):987–1030, 2009a.
- Zhongmin Wang. Station level gasoline demand in an australian market with regular price cycles. *Australian Journal of Agricultural and Resource Economics*, 53(4):467–483, 2009b.

Appendix

Proof of Lemma 1. First note that every $h \in \mathcal{H}^\infty$ is an element of $(\mathcal{Y} \times \mathcal{Y})^\mathbb{N}$.

\mathcal{Y} is finite so it is compact under the discrete topology. $(\mathcal{Y} \times \mathcal{Y})^\mathbb{N}$ is a countable product of compact sets so it is compact by Tychonoff’s theorem.

Now fix any strategy profile s . $\mathcal{H}^\infty(s)$ is a subset of $(\mathcal{Y} \times \mathcal{Y})^\mathbb{N}$.

Now let $C_t(h^t) \subseteq \mathcal{H}^\infty$ be the open cylinder defined by fixing the first $t - 1$ elements of \mathcal{H}^∞ equal to h^t . Then:

$$[\mathcal{H}^\infty(s)]^C = \bigcup_t \bigcup_{h^t \notin \mathcal{H}^t(s)} C_t(h^t)$$

The complement of $\mathcal{H}^\infty(s)$ is therefore a union of open sets. So $\mathcal{H}^\infty(s)$ is closed under the product topology. It is thus compact.

For a fixed $h \in \mathcal{H}^\infty(s)$ define:

$$f_i^T(s, h) = \sum_{t \geq 1}^T \delta^{t-1} u(s_i(h^t), y^t)$$

$$f_i(s, h) = \sum_{t \geq 1}^\infty \delta^{t-1} u(s_i(h^t), y^t)$$

Notice that $f_i^T(s, h)$ is a continuous function of h for every finite T . That is because any preimage of f_i^T is an open cylinder in \mathcal{H}^∞ . Note that I am referring to continuity with respect to the entire sequence of public signals. Furthermore:

$$\left| f_i^T(s, h) - f_i(s, h) \right| \leq M \delta^T$$

$f_i(s, h)$ is continuous because it is the uniform limit of partial sums and each partial sum is continuous.⁸

The minimum of a continuous function on a compact set exists. So $\min_h f_i(s, h)$ exists.

Notice that $\min_h f_i(s, h) = \min_{h \in \mathcal{H}^\infty(s)} \sum_{t \geq 1} \delta^{t-1} u(s_i(h^t), y_i^t, y_j^t)$.

Exactly the same argument can be used for any continuation strategy and any history h^t .

□

Proof of Lemma 2. Suppose the condition in Definition 2 is satisfied. Then for both i and every h^t :

$$\min_{y_j^t \in \mathcal{Y}_j(s_j | h^t)} \left\{ g_i^t(s_i, s_j | h^t, y_j^t) - g_i^t(\tilde{s}_i, s_j | h^t, y_j^t) \right\} \geq 0$$

$$\min_{y_j^t \in \mathcal{Y}_j(s_j | h^t)} g_i^t(s_i, s_j | h^t, y_j^t) - \min_{y_j^t \in \mathcal{Y}_j(s_j | h^t)} g_i^t(\tilde{s}_i, s_j | h^t, y_j^t) \geq 0$$

$$g_i^t(s_i, s_j | h^t) - g_i^t(\tilde{s}_i, s_j | h^t) \geq 0$$

where $\tilde{s}_i \neq s_i$ is a one shot deviation.

⁸The uniform limit of continuous functions is continuous.

Recall that guaranteed payoffs are “continuous at infinity” because of discounting. Moreover, they can be written recursively so the usual one shot deviation principle applies.⁹

That means $g_i^t(s_i, s_j \mid h^t) - g_i^t(\hat{s}_i, s_j \mid h^t) \geq 0$ for an arbitrary $\hat{s}_i \neq s_i$. So for every $y_j^t \in \mathcal{Y}_j(s_j \mid h^t)$:

$$\begin{aligned} g_i^t(\hat{s}_i, s_j \mid h^t, y_j^t) &= \min_{\hat{y}_i^t \in \mathcal{Y}_i(\hat{s}_i \mid h^t)} \left\{ (1 - \delta) u_i(\hat{a}_i^t, \hat{y}_i^t, y_j^t) + \delta g_i^{t+1}(\hat{s}_i, s_j \mid h^t, \hat{y}_i^t, y_j^t) \right\} \\ &\leq \min_{\hat{y}_i^t \in \mathcal{Y}_i(\hat{s}_i \mid h^t)} \left\{ (1 - \delta) u_i(\hat{a}_i^t, \hat{y}_i^t, y_j^t) + \delta g_i^{t+1}(s_i, s_j \mid h^t, \hat{y}_i^t, y_j^t) \right\} \\ &\leq \min_{y_i^t \in \mathcal{Y}_i(s_i \mid h^t)} \left\{ (1 - \delta) u_i(a_i^t, y_i^t, y_j^t) + \delta g_i^{t+1}(s_i, s_j \mid h^t, y_i^t, y_j^t) \right\} \\ &= g_i^t(s_i, s_j \mid h^t, y_j^t) \end{aligned}$$

But that is precisely $r_i(s \mid h^t) \leq 0$.

□

Proof of Lemma 3. Fix any mixed strategy PXE. Let σ be the strategy profile.

Fix any public history h^t .

There must exist some pure strategy s_i such that s_i is in the support of σ_i and:

$$g_j^t(\sigma_i, \sigma_j \mid h^t) = g_j^t(s_i, \sigma_j \mid h^t)$$

Meaning there must be some s_i which is worst for firm j after h^t is realized.

Note that:

$$g_i^t(s_i, \sigma_j \mid h^t) = g_i^t(\sigma_i, \sigma_j \mid h^t)$$

That is because mixing requires firm i to be indifferent between all the strategies in the support of

⁹See Blackwell (1965).

σ_i .

Now suppose that firm i plays the pure action $s_i(h^t)$ after h^t is realized but continues to play according to σ_i after all other histories.

Recall that PXE requires ex post incentive compatibility.

That means firm j is indifferent between all actions in the support of $\sigma_j(h^t)$ even after firm

The fact that σ is a PXE gives us that for all y_i^t in the support of σ_i and all s_j in the support of σ_j :

$$\min_{y_j^t} (1 - \delta) u_j(\sigma_j(h^t), y_j^t, y_i^t) + \delta g_j^{t+1}(\sigma | h^t, y_i^t, y_j^t) = \min_{y_j^t} (1 - \delta) u_j(s_j(h^t), y_j^t, y_i^t) + \delta g_j^{t+1}(\sigma | h^t, y_i^t, y_j^t)$$

Now define a pure strategy s_i^* such that $s_i^*(h^t)$ is the worst action for firm j after every history h^t .

Define s_j^* analogously.

I claim that the strategy profile s^* is a PXE and $g(s^* | h^t) = g(\sigma | h^t)$ for all h^t .

Consider any firm i and public history h^t .

PXE gives:

$$\begin{aligned} g_i^t(\sigma_i, \sigma_j | h^t) &= g_i^t(s_i^*, \sigma_j | h^t) \\ &= g_i^t(s_i^*, s_j^* | h^t) \end{aligned}$$

The second equality comes from the

□

Proof of Proposition 1. Fix any history h^t such that $s_i(h^t) = p_i^t \neq y_i^{t-1}$.

Fix any y_j^t which is realized with positive probability.

Consider a one-shot deviation under which Firm i plays $p_i^t = y_i^{t-1}$.

Firm i 's observed price at the end of period t is certain to be $y_i^t = y_i^{t-1}$.

So the guaranteed payoff after this deviation is:

$$(1 - \delta) u_i(y_i^{t-1}, y_j^t) + \delta g_i^{t+1}(s \mid h^t, y_i^{t-1}, y_j^t)$$

Compare this to the guaranteed payoff in equilibrium:

$$\min_{y_i^t} (1 - \delta) u_i(p_i^t, y_i^t, y_j^t) + \delta g_i^{t+1}(s \mid h^t, y_i^t, y_j^t) \leq (1 - \delta) u_i(p_i^t, y_i^{t-1}, y_j^t) + \delta g_i^{t+1}(s \mid h^t, y_i^{t-1}, y_j^t)$$

PXE requires that the guaranteed payoff in equilibrium is greater than the guaranteed payoff after the deviation:

$$(1 - \delta) u_i(p_i^t, y_i^{t-1}, y_j^t) + \delta g_i^{t+1}(s \mid h^t, y_i^{t-1}, y_j^t) \geq (1 - \delta) u_i(y_i^{t-1}, y_i^{t-1}, y_j^t) + \delta g_i^{t+1}(s \mid h^t, y_i^{t-1}, y_j^t)$$

$$u_i(p_i^t, y_i^{t-1}, y_j^t) \geq u_i(y_i^{t-1}, y_i^{t-1}, y_j^t)$$

This establishes the desired inequality and completes the proof.

□